

Udiddit, a social news aggregator

Introduction

Udiddit, a social news aggregation, web content rating, and discussion website, is currently using a risky and unreliable Postgres database schema to store the forum posts, discussions, and votes made by their users about different topics.

The schema allows posts to be created by registered users on certain topics, and can include a URL or a text content. It also allows registered users to cast an upvote (like) or downvote (dislike) for any forum post that has been created. In addition to this, the schema also allows registered users to add comments on posts.

Here is the DDL used to create the schema:

```
CREATE TABLE bad_posts (  
    id SERIAL PRIMARY KEY,  
    topic VARCHAR(50),  
    username VARCHAR(50),  
    title VARCHAR(150),  
    url VARCHAR(4000) DEFAULT NULL,  
    text_content TEXT DEFAULT NULL,  
    upvotes TEXT,  
    downvotes TEXT  
);  
  
CREATE TABLE bad_comments (  
    id SERIAL PRIMARY KEY,  
    username VARCHAR(50),  
    post_id BIGINT,  
    text_content TEXT  
);
```

Part I: Investigate the existing schema

As a first step, investigate this schema and some of the sample data in the project's SQL workspace. Then, in your own words, outline three (3) specific things that could be improved about this schema. Don't hesitate to outline more if you want to stand out!

1. The post_id in bad_comments should use a foreign key constraint to make sure comments are related to their posts.
2. The post_id should use INTEGER instead of BIGINT.
3. Usernames should be stored into another table to save space.
4. Votes should be stored into another table to save space.

Part II: Create the DDL for your new schema

Having done this initial investigation and assessment, your next goal is to dive deep into the heart of the problem and create a new schema for Udiddit. Your new schema should at least reflect fixes to the shortcomings you pointed to in the previous exercise. To help you create the new schema, a few guidelines are provided to you:

1. Guideline #1: here is a list of features and specifications that Udiddit needs in order to support its website and administrative interface:
 - a. Allow new users to register:
 - i. Each username has to be unique
 - ii. Usernames can be composed of at most 25 characters
 - iii. Usernames can't be empty
 - iv. We won't worry about user passwords for this project
 - b. Allow registered users to create new topics:
 - i. Topic names have to be unique.
 - ii. The topic's name is at most 30 characters
 - iii. The topic's name can't be empty
 - iv. Topics can have an optional description of at most 500 characters.
 - c. Allow registered users to create new posts on existing topics:
 - i. Posts have a required title of at most 100 characters
 - ii. The title of a post can't be empty.
 - iii. Posts should contain either a URL or a text content, **but not both**.
 - iv. If a topic gets deleted, all the posts associated with it should be automatically deleted too.
 - v. If the user who created the post gets deleted, then the post will remain, but it will become dissociated from that user.
 - d. Allow registered users to comment on existing posts:
 - i. A comment's text content can't be empty.
 - ii. Contrary to the current linear comments, the new structure should allow comment threads at arbitrary levels.
 - iii. If a post gets deleted, all comments associated with it should be automatically deleted too.
 - iv. If the user who created the comment gets deleted, then the comment will remain, but it will become dissociated from that user.
 - v. If a comment gets deleted, then all its descendants in the thread structure should be automatically deleted too.
 - e. Make sure that a given user can only vote once on a given post:
 - i. Hint: you can store the (up/down) value of the vote as the values 1 and -1 respectively.
 - ii. If the user who cast a vote gets deleted, then all their votes will remain, but will become dissociated from the user.

- iii. If a post gets deleted, then all the votes for that post should be automatically deleted too.
2. Guideline #2: here is a list of queries that Uddidit needs in order to support its website and administrative interface. Note that you don't need to produce the DQL for those queries: they are only provided to guide the design of your new database schema.
- a. List all users who haven't logged in in the last year.
 - b. List all users who haven't created any post.
 - c. Find a user by their username.
 - d. List all topics that don't have any posts.
 - e. Find a topic by its name.
 - f. List the latest 20 posts for a given topic.
 - g. List the latest 20 posts made by a given user.
 - h. Find all posts that link to a specific URL, for moderation purposes.
 - i. List all the top-level comments (those that don't have a parent comment) for a given post.
 - j. List all the direct children of a parent comment.
 - k. List the latest 20 comments made by a given user.
 - l. Compute the score of a post, defined as the difference between the number of upvotes and the number of downvotes
3. Guideline #3: you'll need to use normalization, various constraints, as well as indexes in your new database schema. You should use named constraints and indexes to make your schema cleaner.
4. Guideline #4: your new database schema will be composed of five (5) tables that should have an auto-incrementing id as their primary key.

Once you've taken the time to think about your new schema, write the DDL for it in the space provided here:

```
CREATE TABLE users (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(25) UNIQUE,  
  CONSTRAINT valid_name CHECK(LENGTH(TRIM(name)) > 0),  
  last_login TIMESTAMP  
);  
  
CREATE INDEX username_index ON users(name);  
  
CREATE TABLE topics (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR(30) UNIQUE,
```

```

        CONSTRAINT valid_name CHECK(LENGTH(TRIM(name)) > 0),
        description VARCHAR(500)
    );

CREATE INDEX topic_name_index ON topics(name);

CREATE TABLE posts (
    id SERIAL PRIMARY KEY,
    title VARCHAR(100),
    CONSTRAINT valid_title CHECK(LENGTH(TRIM(title)) > 0),
    url TEXT,
    content TEXT,
    CONSTRAINT contain_url_or_content CHECK((LENGTH(url) > 0 AND
LENGTH(content) = 0) OR (LENGTH(url) = 0 AND LENGTH(content) > 0)),
    topic_id INTEGER REFERENCES topics ON DELETE CASCADE,
    user_id INTEGER REFERENCES users ON DELETE SET NULL,
    create_time TIMESTAMP
);

CREATE INDEX posts_user_id_index ON posts(user_id);
CREATE INDEX posts_topic_id_index ON posts(topic_id);

CREATE TABLE comments (
    id SERIAL PRIMARY KEY,
    content TEXT,
    CONSTRAINT valid_content CHECK(LENGTH(TRIM(content)) > 0),
    post_id INTEGER REFERENCES posts ON DELETE CASCADE,
    user_id INTEGER REFERENCES users ON DELETE SET NULL,
    top_level BOOLEAN,
    parent_id INTEGER REFERENCES comments ON DELETE CASCADE,
    create_time TIMESTAMP
);

CREATE INDEX comments_user_id_index ON comments(user_id);
CREATE INDEX comments_post_id_index ON comments(post_id);

CREATE TABLE votes (
    value SMALLINT,
    CONSTRAINT valid_value CHECK(value = -1 OR value = 1),
    post_id INTEGER REFERENCES posts ON DELETE CASCADE,
    user_id INTEGER REFERENCES users ON DELETE SET NULL,
    id SERIAL PRIMARY KEY
);

CREATE INDEX votes_user_id_index ON votes(user_id);
CREATE INDEX votes_post_id_index ON votes(post_id);

```

Part III: Migrate the provided data

Now that your new schema is created, it's time to migrate the data from the provided schema in the project's SQL Workspace to your own schema. This will allow you to review some DML and DQL concepts, as you'll be using INSERT...SELECT queries to do so. Here are a few guidelines to help you in this process:

1. Topic descriptions can all be empty
2. Since the bad_comments table doesn't have the threading feature, you can migrate all comments as top-level comments, i.e. without a parent
3. You can use the Postgres string function **regexp_split_to_table** to unwind the comma-separated votes values into separate rows
4. Don't forget that some users only vote or comment, and haven't created any posts. You'll have to create those users too.
5. The order of your migrations matter! For example, since posts depend on users and topics, you'll have to migrate the latter first.
6. Tip: You can start by running only SELECTs to fine-tune your queries, and use a LIMIT to avoid large data sets. Once you know you have the correct query, you can then run your full INSERT...SELECT query.
7. **NOTE:** The data in your SQL Workspace contains thousands of posts and comments. The DML queries may take at least 10-15 seconds to run.

Write the DML to migrate the current data in bad_posts and bad_comments to your new database schema:

```
INSERT INTO topics(name)
SELECT DISTINCT topic
FROM bad_posts;

INSERT INTO users(name)
SELECT DISTINCT username
FROM bad_posts;

INSERT INTO users(name)
SELECT DISTINCT username
FROM bad_comments bc
WHERE NOT EXISTS (
    SELECT name
    FROM users u
    WHERE u.name = bc.username
);
```

```

INSERT INTO users(name)
SELECT DISTINCT username
FROM (
    SELECT regexp_split_to_table(upvotes, ',') username
    FROM bad_posts) upvote_users
WHERE NOT EXISTS(
    SELECT name
    FROM users u
    WHERE u.name = upvote_users.username
);

INSERT INTO users(name)
SELECT DISTINCT username
FROM (
    SELECT regexp_split_to_table(downvotes, ',') username
    FROM bad_posts) downvote_users
WHERE NOT EXISTS(
    SELECT name
    FROM users u
    WHERE u.name = downvote_users.username
);

INSERT INTO posts(id, title, url, content, topic_id, user_id)
SELECT bp.id, LEFT(bp.title, 100), bp.url, bp.text_content, t.id,
u.id
FROM bad_posts bp
JOIN users u
ON u.name = bp.username
JOIN topics t
ON t.name = bp.topic;

INSERT INTO comments(content, post_id, user_id, top_level)
SELECT bc.text_content, p.id, u.id, true
FROM bad_comments bc
JOIN users u
ON u.name = bc.username
JOIN posts p
ON p.id = bc.post_id;

INSERT INTO votes(value, post_id, user_id)
WITH up AS (
    SELECT regexp_split_to_table(upvotes, ',') username, id
    FROM bad_posts
)
SELECT 1, up.id, u.id
FROM users u
JOIN up

```

```
ON up.username = u.name;

INSERT INTO votes(value, post_id, user_id)
WITH down AS (
    SELECT regexp_split_to_table(downvotes, ',') username, id
    FROM bad_posts
)
SELECT -1, down.id, u.id
FROM users u
JOIN down
ON down.username = u.name;
```